

ALBERT-LUDWIGS-UNIVERSITÄT  
FREIBURG

# A Neural Network for RPCA

A PROJECT IN  
STOCHASTIC MACHINE LEARNING

*Eric Brunner, Aron Distelzweig, Mirko Grimm, Bastian  
Schnitzer, Simon Beyer*

January 31, 2021

# Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>Denise</b>	<b>2</b>
<b>3</b>	<b>Generation of training data</b>	<b>3</b>
<b>4</b>	<b>Testing and Comparison</b>	<b>4</b>
4.1	The Minimization Problem solved by PCP . . . . .	4
4.2	PCP Algorithm . . . . .	5
4.3	Evaluation of trained Denise on Financial Data . . . . .	6
4.4	Portfolio Correlations from DAX 30 with PCP and Denise . . . . .	6
4.5	Correlation matrix of personality features with PCP and Denise . . . . .	7
<b>5</b>	<b>Outlook</b>	<b>11</b>
5.1	Convolutional Network . . . . .	11
5.2	RPCA for arbitrary matrices . . . . .	11
	<b>References</b>	<b>15</b>

# 1 Introduction

In Principal Component Analysis (PCA) we aim at finding the principal components of a set of data points. The principal components of a set of points in  $\mathbb{R}^n$  are a sequence of  $n$  vectors. They are recursively defined as the  $i^{th}$  vector being the direction of a line that best fits the data while being orthogonal to the first  $i - 1$  vectors. The line is obtained through minimizing its average squared distance from the data points. Intuitively, one can think of PCA as fitting a  $p$ -dimensional ellipsoid to the data, where each axis of the ellipsoid represents a principal component. Mathematically this corresponds to the eigenvectors of the data's covariance matrix. Hence, the principal components form an orthonormal basis of the space  $\mathbb{R}^n$ . PCA is heavily used as an exploratory tool in data analysis. One big application is in quantitative finance (or generally time series analysis), where one is interested in the principal components of the empirical covariance matrix of certain financial assets. Focusing only on the largest principal components, which still hold the most important information about the data, leads to a dimensional reduction of the large dataset such that an efficient analysis of the data is still feasible.

In this project we specifically look at a modification of PCA, which is the Robust Principal Component Analysis. Here our aim is to recover a low rank Matrix  $L$  from, possibly highly, corrupted matrices  $M$ , in the sense that some matrix entries are faulty, for example through imprecise measurements. More precisely, we start with a data matrix (or its covariance matrix)  $M \in \mathbb{R}^{m \times n}$  with corrupted entries and attempt to find a decomposition into a sum  $M = L + S$ .  $L$  is a matrix of low rank while the corrupted entries are filtered out into a sparse matrix  $S$  (which means that a lot of entries are zero). In this project we will restrict to the RPCA of symmetric positive semi-definite matrices  $M$ , which are for example given as empirical covariance matrices.

The usual state-of-the-art algorithms for (R)PCA are computationally demanding and, hence, impractical in some applications like finance, where instantaneous calculation might be required. Therefore, the use of a neuronal network is seen to yield a valuable alternative for the design of an efficient decomposition tool. In this project our aim is to implement the algorithm "Denise" (see [1]) that aims at solving the robust PCA for semi-definite matrices through direct learning of the decomposition map  $M \mapsto L + S$  via a deep neural network. We train Denise on a randomly generated synthetic dataset and evaluate the performance of the deep neural network on synthetic and real-world covariance matrices. The advantage of a neural network compared to the standard methods is that once trained the neural network delivers a function that outputs the desired decomposition of a dataset instantaneously, where using standard methods like convex optimization are computationally expensive. The algorithms are matrix specific (meaning that one has to use the algorithm for every new data), which is, therefore, very slow especially when data becomes large. For example in finance, one needs robust low rank estimation of covariance matrices of hundreds of assets instantaneously where the high performance of a neural network is very practical. As argued in [1], the results achieved by Denise are comparable to several state-of-the-art algorithms in terms of decomposition quality but outperforms all existing algorithms by computation time. As explained above, this is achieved by learning a single evaluation function that takes a matrix as an input and outputs the desired decomposition.

In section 2 the considered algorithm is explained in detail. Especially we introduce the objective function on which we train the neural network and describe the network architecture. Subsequently in section 4 we introduce the principal component pursuit (PCP) algorithm, which will be used as benchmark to assess the performance of the neural networks approach. We evaluate PCP as well as the trained Denise network on a Portfolio correlation matrix of five exemplary stocks in the DAX 30 and compare the resulting decompositions. Finally, in section 5, we summarize approaches on how to develop our project further, and which additional tasks we would like to carry out in the course of this semester.

## 2 Denise

We consider  $\mathbb{S}_n \subseteq \mathbb{R}^{n \times n}$  to be the set of  $n$ -by- $n$  symmetric matrices and  $\mathbb{P}_n \subseteq \mathbb{S}_n$  to be the subset of positive semi-definite matrices and  $\mathbb{P}_{k,n} \subseteq \mathbb{P}_n$  the subset of matrices with rank at most  $k$ . As the input of our deep neuronal network we consider a Matrix  $M = (M_{ij})_{ij} \in \mathbb{P}_n$ .  $M$  is to be decomposed as a sum  $M = L + S$  where  $L = (L_{ij})_{ij} \in \mathbb{P}_{k,n}$  is of rank at most  $k$  and a sparse matrix  $S = (S_{ij})_{ij} \in \mathbb{P}_n$ . Since  $L$  is assumed to be symmetric, by the Cholesky decomposition, we can represent it as  $L = UU^T$ , where  $U = (U_{ij})_{ij} \in \mathbb{R}^{n \times k}$ . Therefore  $M$  can be expressed as  $M = UU^T + S$ .

**Loss function** We input  $M$  in the neural network, which should faithfully output the low rank matrix  $U$ . Hence, we want to minimize the difference between  $M$  and  $UU^T$ , which is equal to  $S$ . A convenient choice of loss-function for the considered neural network is, therefore, given by the  $l_1$ -matrix-norm of  $S$

$$\|UU^T - M\|_{l_1} = \|S\|_{l_1} \quad (2.1)$$

The  $l_1$ -norm is a common choice to guarantee sparsity of  $S$ .

**Architecture** As the matrix  $M$  is symmetric, we can reduce the input from  $n^2$  to  $n(n+1)/2$  by taking the triangular lower matrix of  $M$ . The lower matrix is then transformed into a vector using the operator  $h$ :

$$h : \mathbb{S}_n \rightarrow \mathbb{R}^{n(n+1)/2}, M \mapsto (M_{1,1}, M_{2,1}, M_{2,2}, \dots, M_{n,1}, \dots, M_{n,n})^T \quad (2.2)$$

Similarly we convert the output vector of the neural network into a matrix with the operator  $g$  defined as

$$g : \mathbb{R}^{n(n+1)/2} \rightarrow \mathbb{R}^{n \times k}, X \mapsto \begin{pmatrix} X_1 & \cdots & X_k \\ \vdots & & \vdots \\ X_{(n-1)k+1} & \cdots & X_{(n-1)k+k} \end{pmatrix} \quad (2.3)$$

Besides the output layer, our multi-layer feed-forward neural network  $\mathcal{N} : \mathbb{R}^{n(n+1)/2} \rightarrow \mathbb{R}^{nk}$  has three hidden dense layers, each exhibiting ReLU-activation function and  $n/2$  nodes.

Using  $h$  and  $g$  the matrix  $U$  is the output of the neural network  $U = g(\mathcal{N}(h(M)))$  and we get the desired matrix  $L = \rho(\mathcal{N}(h(M)))$  for

$$\rho : \mathbb{R}^{rd} \rightarrow \mathbb{P}_{r,d}, X \mapsto g(X)g(X)^T \quad (2.4)$$

### 3 Generation of training data

For the training of Denise a synthetic dataset, consisting of randomly generated matrices with appropriate decomposition properties, is created. In detail, we construct a sample of positive semidefinite  $n$ -by- $n$  matrices  $M$  that can be decomposed as

$$M = L_0 + S_0 \quad (3.1)$$

where  $L_0$  is a known matrix of rank  $k_0 \leq n$  and  $S_0$  a known matrix with a given sparsity  $s_0$ . Here, we understand a *sparse matrix* as a matrix containing a lot of zeros and define the ratio between the number of zero-valued entries and the total number of entries as its *sparsity*. In the process of data generation, we follow a reverse approach by constructing first the matrices  $L_0$  and  $S_0$  with the required properties and merge it together to a matrix  $M$  which has by definition the desired decomposition.

For the construction of the low-rank matrix  $L_0$ , we collect  $nk_0$  samples of independent standard normal random variables into an  $n$ -by- $k_0$  matrix  $U$  and set  $L_0 = UU^T$ . This construction scheme guarantees symmetry and positive semidefiniteness of  $L_0$  as well as rank  $L_0 \leq k_0$ .

To construct the symmetric positive semidefinite sparse matrix  $S_0$ , we first take a sample of a uniformly random pair  $(i, j)$  with  $1 \leq i < j \leq n$  that defines four non-zero entries of an  $n$ -by- $n$  matrix  $\tilde{S}_0$ . The off-diagonal elements  $(i, j)$  and  $(j, i)$  are set to value  $b$  as the realization of a uniformly random variable in  $[-1, 1]$  while the diagonal elements  $(i, i)$  and  $(j, j)$  are set to value  $a$  as the realization of a uniformly random variable in  $[|b|, 1]$ . Due to its construction  $\tilde{S}_0$  is positive semidefinite. Finally, we receive the matrix  $S_0$  by summing the realizations  $\tilde{S}_0^{(i,j)}$  over the pairs  $(i, j)$  until the pursued sparsity is reached.

According to the described approach, we define a class *SyntheticMatrixSet* which is, by its methods, able to create a randomly generated matrix  $M$  together with  $L_0$  and  $S_0$  as both parts of its decomposition  $M = L_0 + S_0$  based on the userspecified arguments: dimension  $n$ , rank  $k_0$  and sparsity  $s_0$ . This class enables the generation of highly extensive datasets for various training purposes.

## 4 Testing and Comparison

After the neural network has been trained with the synthetic data, as described in section 2, it can be tested on synthetic data, where the decomposition is known by construction, as well as on real world data. To compare the decompositions of real world data, another algorithm is needed for benchmarking, which calculates the decomposition of an arbitrary positive semidefinite matrix  $M$  into a low rank, symmetric and positive-semidefinite matrix  $L_0$  plus a sparse matrix  $S_0$ :

$$M = L_0 + S_0. \quad (4.1)$$

With *Principal Component Pursuit* (PCP), there exists an algorithm, which, under some suitable assumptions, calculates the decomposition exactly via singular value decomposition (SVD) [2]. The assumptions and the main ideas of this algorithm is presented in the first subsection. In the second subsection the results of the decomposition of portfolio correlation matrices via the AI algorithm Denise with the results of the Principal Component Pursuit algorithm.

### 4.1 The Minimization Problem solved by PCP

Let  $M$  be a given element of  $\mathbb{R}^{n_1 \times n_2}$ .  $\|\cdot\|_*$  denotes the nuclear norm, i.e. the sum over the singular values of a matrix  $\|M\|_* := \sum_i \sigma_i(M)$ .  $\|\cdot\|_1$  is the well known  $\ell_1$  norm  $\|M\|_1 = \sum_{ij} |M_{ij}|$ . The PCP algorithm solves the convex optimization problem

$$\text{minimize} \quad \|L\|_* + \|S\|_1, \quad \text{where } L + S = M \quad (4.2)$$

exactly, if the the low-rank component  $L_0$  fulfills a "incoherence" condition, and that the sparse component is "reasonably sparse". The meaning of this "incoherence" condition for  $L_0$  and the "reasonable" sparsity of  $S_0$  is explained in [2, subsection 1.3]. We summarize the main points real for quadratic matrices:

- (i) Let  $U\Sigma V^\top$  the singular singular value decomposition of  $L_0 \in \mathbb{R}^{n \times n}$  with rank  $k \geq n$ , i.e.

$$L_0 = U\Sigma V^\top = \sum_{i=1}^k \sigma_i u_i v_i^\top, \quad (4.3)$$

where  $U = (u_1, \dots, u_k), V = (v_1, \dots, v_k) \in O(n)$ ,  $\Sigma = \text{diag}(\sigma_1, \dots, \sigma_k, 0, \dots, 0) \in \mathbb{R}^{n \times n}$ .  $\sigma_1, \dots, \sigma_k$  are the singular values and  $u_i$  and  $v_i$ ,  $i = 1, \dots, k$ , are the left-singular and right-singular vectors for  $\sigma_i$ , respectively. Then the matrix  $L_0$  is called incoherent, with parameter  $\mu$ , if

$$\max_i \|U e_i\|^2 \geq \frac{\mu k}{n^2}, \quad \max_i \|V e_i\|^2 \geq \frac{\mu k}{n^2}, \quad \|UV^\top\|_\infty \geq \frac{\sqrt{\mu k}}{n}. \quad (4.4)$$

$e_i$  are the canonical basis vectors of  $\mathbb{R}^n$ .

- (ii) The positions of the nonzero elements of the sparsity matrix are selected uniformly random.

If (i) is fulfilled, the matrix  $L_0$  is considered as not sparse. With (ii) we try to prevent, that the nonzero elements are only in one, or few columns of the sparsity matrix. For example if the entries of  $S_0$  except the first column are all zero, and the first column of  $S_0$  is the negative of the first column of  $L_0$ , then it is impossible to recover the low rank component and sparse component exactly. To avoid, such variety of possibilities for the decomposition (ii) is a reasonable assumption.

## 4.2 PCP Algorithm

In this subsection, a brief description of the PCP algorithm, which we use for comparison with Denise, is given. There are different strategies to solve the problem (4.2) numerically. As described in [2], we consider an *augmented Lagrange multiplier*. This is why Candes et al. named this method the ALM method. (4.2) is equivalent to the minimization of the following *augmented Lagrangian*

$$\mathcal{L}(L, S, Y) = \|L\|_* + \lambda \|S\|_1 + \langle Y, M - L - S \rangle + \frac{\mu}{2} \|M - L - S\|_F^2. \quad (4.5)$$

Here  $\langle \cdot, \cdot \rangle$  is defined as  $\langle A, B \rangle = \text{tr}(A^\top B)$ , with real quadratic matrices  $A, B$ .  $\|\cdot\|_F$  is the Frobenius norm. One can show, that

$$\arg \min_S \mathcal{L}(L, S, Y) = \mathcal{S}_{\lambda\mu}(M - L + \mu^{-1}Y), \quad (4.6)$$

$$\arg \min_L \mathcal{L}(L, S, Y) = \mathcal{D}_\mu(M - S - \mu^{-1}Y), \quad (4.7)$$

where  $\mathcal{S}_\tau : \mathbb{R}^{n \times n} \rightarrow \mathbb{R}^{n \times n} : (X_{ij})_{ij} \mapsto (\text{sgn}(X_{ij}) \max(|X_{ij}| - \tau, 0))_{ij}$ , is the extension of the shrinkage operator in  $\mathbb{R}$  to  $\mathbb{R}^{n \times n}$ .  $\mathcal{D}_\tau(X)$  is defined as  $\mathcal{D}_\tau(X) = U\mathcal{S}_\tau(\Sigma)V^\top$ , where  $U\Sigma V^\top$  is the SVD of  $X$ . Hence, the following algorithm, taken from [2, p. 29], is productive

1. **Initialize:**  $S_0 = Y_0 = 0, \mu > 0$ .

2. **While** not converged **do**

$$L_{k+1} = \mathcal{D}_\mu(M - S_k - \mu^{-1}Y_k) \quad (4.8a)$$

$$S_{k+1} = \mathcal{S}_{\lambda\mu}(M - L_{k+1} + \mu^{-1}Y_k) \quad (4.8b)$$

$$Y_{k+1} = Y_k + \mu(M - L_{k+1} - S_{k+1}) \quad (4.8c)$$

3. **Return:**  $L, S$ .

With the calculations of the second step, it is avoided to solve a sequence of convex programs. To achieve good relative accuracy, only a few iteration steps are necessary [2, section 3].

### 4.3 Evaluation of trained Denise on Financial Data

To get a first impression of Denise performance on real-world data, we apply Denise on a set of 10-by-10 correlation matrices of ten stocks out of DAX 30, namely Allianz, BASF, Bayer, Beiersdorf, BMW, Continental, Siemens, Merck, Daimler, VW. As a data basis, we exploit the price data of the aforementioned stocks from the last 2 years, starting from January 15th in 2021. The empirical correlation matrices are determined retrospectively with an offset of one day, while the correlations in time step  $n$  are calculated based on closing prices within the last 35 days. In the end, we obtain 472 correlation matrices  $\mathbf{R}^{(n)}, n = 1, 2, \dots, 472$  for further numerical tests by this approach.

The stock prices required to calculate the empirical correlation matrices are downloaded from *Yahoo! Finance* and then transferred into a CSV file. This CSV file is imported into *Jupyter Notebook* and the price data is stored in vectors  $X^{(i)}, i = 1, \dots, 10$ , beginning with the latest value, in order to calculate the empirical correlation matrices  $\mathbf{R}^{(n)}, n = 1, 2, \dots, 472$  according to

$$\mathbf{R}_{i,j}^{(n)} = \frac{\sum_{k=n}^{n+N-1} (X_k^{(i)} - \bar{X}^{(i)})(X_k^{(j)} - \bar{X}^{(j)})}{\sqrt{\sum_{k=n}^{n+N-1} (X_k^{(i)} - \bar{X}^{(i)})^2 \sum_{k=n}^{n+N-1} (X_k^{(j)} - \bar{X}^{(j)})^2}}, \quad (4.9)$$

where  $X_k^{(i)}$  is the price of the  $i$ -th stock at the end of the  $k$ -th trading day and  $N$  the total number of analyzed trading days, within the relevant time period.

### 4.4 Portfolio Correlations from DAX 30 with PCP and Denise

For test purposes we applied the PCP algorithm on 472 empirical correlation matrices of the ten DAX 30 companies, mentioned before. Due to our approach in calculating the different correlation matrices, their structure changes quite slowly. Hence, we concentrate our investigations only on 10 matrices with equidistant time steps.

We train the neural network on a synthetic data set of 1000000  $10 \times 10$ -matrices of rank three. Training and validation loss are illustrated in Fig. 1. The comparison of the RPCA decomposition of the correlations matrices into rank 3 matrix  $L$  and sparse  $S$  for both methods is shown in Fig. 2. Representatively, we have chosen only one decomposition, because the color plots of the ten matrices are quite similar. This similarity allows us to restrict further investigations concerning the attributes of the decompositions on only one single matrix. As one can see, the decomposition obtained from our neural network suffers significant outliers on the diagonal elements in  $L$  as well as  $S$ , which is not apparent in the PCP decomposition. One explanation for the occurrence of the outliers on the diagonal can be the structure of the training matrices. The normalization of the correlation matrices causes, that the diagonal elements are all set to 1. Since this particular property is missing in the training data, the network could not process the diagonal elements accurately. Due to this outlier not much structure is visible in  $L$  obtained from the neural network. However, apart from the diagonal elements the  $S$ -matrices of both methods seem to agree to some approximation (respecting the corresponding color code). As a






Figure 1: Training and validation loss of the network training for finance data.

metric, to measure how well the two algorithms agree, we determine the relative errors  $\epsilon_{\text{rel}}(S_{\text{PCP}}, S_{\text{NN}})$ ,  $\epsilon_{\text{rel}}(L_{\text{PCP}}, L_{\text{NN}})$  as the relative  $l_2$ -distance

$$\epsilon_{\text{rel}}(A, B) = \frac{\|A - B\|_{l_2}}{\|A\|_{l_2}} \quad (4.10)$$

between  $S_{\text{PCP}}, S_{\text{NN}}$ , respectively between  $L_{\text{PCP}}, L_{\text{NN}}$  (the PCP-, respectively neural-network-prediction for  $S$  and  $L$ ):

$$\epsilon_{\text{rel}}(S_{\text{PCP}}, S_{\text{NN}}) = 5,788, \quad \epsilon_{\text{rel}}(L_{\text{PCP}}, L_{\text{NN}}) = 1,098. \quad (4.11)$$

## 4.5 Correlation matrix of personality features with PCP and Denise

Here we compare the performance of the RPCA of our neural network approach and the benchmark PCP algorithm on 25 personality self report items obtained from approximately 2800 individuals. The studied 25 personality features are grouped in the five categories agreeableness, conscientiousness, extraversion, neuroticism, and openness, see [personality-project]. The aim of PCA is to recover these five underlying putative factors from the data.

In this subsection we compare the performance of PCP and our neural network approach to obtain the RPCA-decomposition of the correlation matrix of the 25 personality features. The correlations are defined as the normalized empirical covariances between the variables  $x_i$

$$M_{ij} = \text{Corr}(x_i, x_j) = \frac{\text{Cov}(x_i, x_j)}{\sqrt{\text{Var}(x_i)\text{Var}(x_j)}}.$$




Figure 2: Comparison of RPCA on the empirical  $10 \times 10$  covariance matrix based on the price development of 10 share certificates (see main text) from the DAX 30. The left panel shows the resulting decomposition  $M = L + S$  obtained from the neural network approach, while the right panel shows the results from PCP.




Figure 3: Training and validation loss of the training of our neural network for the personality data. Curiously, the validation loss is always smaller than the training loss.

As in the previous subsection, we calculate the RPCA-decomposition  $M = L + S$  by means of our neural network approach as well as by using the PCP method introduced in 4.2. The obtained low rank matrix  $L$  and the sparse part  $S$  containing corruptions of the input matrix are shown in Fig. 4. The low-rank ( $k = 10$ ) is determined by the PCP-algorithm, and subsequently used in the definition of the neural network. The training performance of the neural network (training and validation loss) is depicted in Fig. 3. The training was performed on 80000 synthetic  $25 \times 25$  positive semi-definite matrices of rank 10.

As in the analysis of the financial data, the neural network decomposition exhibits very few strong outliers on the diagonal, which suppress the visual structure of the obtained matrices  $L, S$  to some extent. As before, these outliers are not present in the PCP results. Hence, the PCP-method seems to function more stable than the network ansatz. We quantify the discrepancy of the resulting decomposition of both approaches by their relative distances

$$\epsilon_{\text{rel}}(S_{\text{PCP}}, S_{\text{NN}}) = 3.24, \quad \epsilon_{\text{rel}}(L_{\text{PCP}}, L_{\text{NN}}) = 2.14. \quad (4.12)$$

The comparison of both methods, in combination with the striking results reported in [1], may be explained by a sub-optimal choice of network-architecture (to few nodes per hidden layer) of our network. Another source of instability might result from the small training sets or not ideally chosen hyper-parameters such as batch-size. This will be analyzed in detail in future work. Moreover, we will analyze, how we can use RPCA to obtain characteristic information about the studied data. For example, it will be interesting to investigate in which way we can recover the five underlying personality categories agreeableness, conscientiousness, extraversion, neuroticism, and openness from the principal components.




Figure 4: Comparison of RPCA of the empirical  $25 \times 25$  covariance matrix of 27 personality features of roughly 2800 individuals (see main text). The left panel shows the resulting decomposition  $M = L + S$  obtained from the neural network approach, while the right panel shows the results from PCP.

## 5 Outlook

In this section we give a brief outlook on some ideas which we like develop further in the course of this project.

### 5.1 Convolutional Network

It is clear that the described neural network ansatz (see section 2), once trained yields a very fast and robust methodology to solve the RPCA problem, compared to state-of-the-art algorithms (as for example the PCP discussed in 4.2). However, the training and the generation of training data can be computationally quite demanding for large matrices. On the one hand, larger matrices require much larger training data sets for obtaining reasonable generalization properties. On the other hand, the necessary network parameters in the dense network implementation will also scale up. To circumvent this scaling problem (at least on the training side) we would like to compare the current dense network topology to a more sparse one, building on convolutionary layers. These architectures, typically, exhibit much fewer trainable parameters than dense networks.

Convolutionary networks are usually used in image recognition. In such scenarios one typically expects strong correlations within the data (pixels of the images), since usually some red pixels will be most probable surrounded by further red ones. By exploiting these correlations, a convolutionary network (with much fewer parameters than a dense one) is often able to perform equally well while allowing for much faster training. Although the considered corrupted symmetric positive semi-definite matrices at input might not boast a similar structure than images, the positive definite character also will lead to correlations between the individual matrix entries. Therefore, it seems reasonable to test a sparser convolutionary network architecture also in this setting.

### 5.2 RPCA for arbitrary matrices

As described in section 2, the RPCA of a symmetric positive semi-definite matrix  $M$  is given by decomposing

$$M = UU^T + S$$

with a sparse matrix  $S$  and a rank  $k$  matrix  $U \in \mathbb{R}^{n \times k}$ . It is clear that this approach not directly generalizes to arbitrary matrices such as data matrices, since  $UU^T$  is by definition positive semi-definite. However, any (also non-quadratic) rank  $k$  matrix  $M \in \mathbb{R}^{n \times m}$  can be decomposed as  $M = UV^T$  with  $U \in \mathbb{R}^{n \times k}$  and  $V \in \mathbb{R}^{m \times k}$ . Therefore, the RPCA problem of arbitrary  $M$  is given by finding a sparse matrix  $S$  and rank  $k$  matrices  $U, V$  as above that fulfill

$$M = UV^T + S.$$

A first approach to this problem, based on neural networks, is to employ two collaborating networks, each similar to the network used so far. Each networks task is to learn the

mapping from input matrix  $M$  to low-rank  $k$  matrix  $U$ , respectively  $V$ , by minimizing the  $l_1$ -distance  $\|M - UV^T\|_{l_1}$  to ensure (as before) sparsity of  $S$ . The training of both networks can be accomplished in an alternating manner. E.g. we first initialize  $V$  at random and train the first network to approximate  $U$ , while in a second step we fix the first network and its output  $U$  to train the second network on the same data to output  $V$ . Based on this (ideally better choice of  $V$ ) we can now repeat this procedure up to a point where both networks converged up to a reasonable degree and stabilize. The collaborative training of both network will naturally take approximately two times as long as the single network employed so far and exhibits, hence, similar computational complexity.




Figure 5: Comparison of RPCA of the empirical  $5 \times 5$  covariance matrix of the prices of five share certificates (see main text) from the DAX 30. The left panel shows the resulting decomposition  $M = L + S$  obtained from the neural network approach, while the right panel shows the results from PCP.




Figure 6: Comparison of RPCA of the empirical  $25 \times 25$  covariance matrix of 25 personality features of roughly 2800 individuals (see main text). The left panel shows the resulting decomposition  $M = L + S$  obtained from the neural network approach, while the right panel shows the results from PCP.



## References

- [1] Calypso Herrera et al. *Denise: Deep Learning based Robust PCA for Positive Semidefinite Matrices*. 2020. arXiv: 2004.13612 [stat.ML].
- [2] Emmanuel J. Candes et al. *Robust Principal Component Analysis?* 2009. arXiv: 0912.3599 [cs.IT].